

A Project Plan for Comparing Statistical and Machine Learning Models for Electricity Demand Forecasting in NSW

Group A:

Kelly Hamilton (z552890)

Daniel Cao (z5267347)

Swapnil Shinde (z5118116)

Faiza Mumtaz (z5657874)

Xavier Sun (z5647547)

Arman Hajisafi (z3542724)

08/09/2026

Page Count: 3 (Excluding cover page, references and appendix)

1. Introduction and Motivation	2
2. Brief Literature Review	2
3. Methods, Software and Data Description	3
4. Activities and Schedule	4
5. References	5
6. Appendix	7

1. Introduction and Motivation

The National Electricity Market (NEM) coordinates electricity generation, trading and distribution across eastern and southern Australia, matching electricity supply with demand in near real time [1]. As electricity cannot be stored at scale, the NEM depends on electricity forecasts to ensure that supply and demand are balanced [2]. Accurate forecasting supports Australia's progress toward international sustainability commitments, such as the United Nations Sustainable Development Goal 7, an area in which Australia has been assessed as "off track" [3]. Additionally, this dependence on accurate forecasting carries financial incentives. For example, generators may rely on these forecasts to decide whether to keep their plants running for further demand or shut down to save costs, however risking the added expense of restarting if demand rises unexpectedly [4].

Because of these real financial consequences and broader sustainability commitments, organisations seek demand analysis to make informed planning and investment decisions. This project is based on a brief from Endgame Economics, a consultancy which provides this kind of analysis. Their clients are typically market bodies that rely on accurate demand forecasts to inform decisions such as policy and investments [5]. Therefore, in this project we propose the following research question; "Which forecasting models perform best for medium to long term electricity demand in New South Wales when evaluated across seasonal variation, temperature changes, and differing demand levels?"

Answering this question directly supports the kind of decision making our potential client may rely on, helping market bodies identify which modelling approaches are most robust under varying conditions, and in turn informing more confident policy and investment decisions. To answer this question, we intend to apply a combination of statistical and machine learning techniques, evaluating each against the seasonal, temperature, and demand level conditions. By using statistical methods, we intend to capture the seasonality of electricity demand. While using machine learning methods, we intend to capture the non-linear, temperature driven variability in demand that time series models alone may not fully capture. By the end of this project, we aim to deliver a comparative evaluation of these models, leveraging existing literature and identifying which best balances accuracy and interpretability across NSW's seasonal and temperature conditions.

2. Brief Literature Review

National Energy Retail Law (NSW) requires companies to provide a reliable supply of energy; achieve targets set by the Australian Energy Market Commission (AEMC); and reduce Australia's greenhouse gases [6]. To ensure sustainability and reliability, efficient and cost-effective prediction of energy demand is imperative for energy retailers [7, 8]. Previous research suggests four possible approaches to meet the needs of retailers considering dynamic city-level data; Time Series Models; Regression Models; Convolutional Neural Networks; and Hybrid Models.

Daily load half hour readings from the Transmission Company of Nigeria (Aug-Dec, 2021) accurately predicted by LSTM, indicating successful short-term prediction, with similar time-step parameters to the current situation [14, A4]. Likewise, in comparison to XGB and ARIMA, LSTM was most accurate; predicting South African energy consumption [13, A3]. Applied to Shanghai office building energy consumption data (24th June- 15 Sept, 2013), LSTM outpaced Backpropagation Network (BP); SVM; and ANN[15, A5]. Further LSTM accuracy obtained via optimisation for timestep [15, A6]. Comparison of U.S. energy consumption trends revealed LR marginally outperformed RF and XGB [11]. XGB trailed RF and LR on Accuracy; precision exposed LR's outperformance [11, A2]. Demonstrating importance of tuning hyperparameters, optimised RF outperformed all compared models (ARIMA; ETS; Prophet; MLP; SVM; ANFIS; GRNN; LSTM; RandNN; XGB, LightGBM; FNM, Nadaraya-Watson

Estimator; MTGNN and hybrid ES-adRNNs) [12]. Renewable energy focused city-level dynamic datasets (Vaasa, Kalmar, Brunei, Thailand) included similar features: Barometric pressure, temperature, humidity, sun radiation, holidays to make short-term load predictions with 1hr and 30min timestep data [10, 16-18]. Vaasa study revealed XGBoost (XGB) the most reliable model over Multiple Linear Regression (MLR); while in Brunei, comparison between ARIMA, SARIMA and hybrid Holt-Winters+SARIMA gave SARIMA the highest recommendation [10, 16, A1, A7]. To forecast Thailand's city region, hybrid CLTSM-RNN produced fortified predictive power against LSTM; RNN and CNN [18]. Hybrid models reveal requirement to balance accuracy with computational cost [16]. In Kalmar, XGB obtained highest error in training, while SARIMAX performed consistently mid-range for both training and testing [17, A8, A9]. FBP outperformed SARIMAX in testing; however, all models produced lower predictive performance than CNN for both training and testing [17]. Sun radiation contributes to prediction accuracy due to consumer solar panels [11, 16, 17]. Holidays apply pressure on demand [17].

Prediction Accuracy is associated with granularity of data [14,17,18]. Given current dataset has a high granularity of 30 min timesteps, modelling choice pursues balance between prediction accuracy and computational cost; providing energy companies with efficient and economically sound forecasts [6-7, 16]. Due to suburb-scope of Bankstown, NSW, there is no spatial variation to consider; therefore, it is suggested to focus on time series and regression methods, reducing resource requirements [12, 15, 17, 18]. Alternatively, exploratory data analysis may investigate increased timesteps to reduce model load and determine impact on accuracy [15]. A leading limitation within literature appears to be the absence of temperature fluctuations consistent with Australian specific-location standards. [A10, 19-25]. The current temperature dataset obtained from Bankstown, NSW contains a much higher temperature range (Max: 44.7°C; Min:-1.3°C; Range: 46°C), presenting an avenue to explore the prediction capabilities of SARIMA, Prophet, Random Forest, XGBoost and LightGBM within broader fluctuation bounds. Due to the nature of energy demand, research question focus and consistent incorporation of multivariate datasets across the research, it is recommended to obtain further data such as school holidays, humidity, sun radiation and/or cloud cover corresponding to decreased solar panel capability and therefore higher demand [10, 11].

3. Methods, Software and Data Description

The project will use a combination of statistical and machine learning methods to forecast electricity demand in New South Wales. This approach reflects findings in the literature that no single modelling technique is sufficient for capturing both the regular seasonal structure of demand and the strong temperature driven variability observed in the NEM [1,7]. The modelling strategy begins with simple, well-established benchmarks and then evaluates more flexible models against them, ensuring that the final choices are supported by evidence rather than complexity alone.

This project will model half-hourly New South Wales electricity demand using a combined statistical and machine learning approach to capture both regular seasonal structures and extreme temperature driven variability [1,7]. To systematically improve upon baseline operator forecasts [1], classic time series methods will be evaluated against flexible machine learning algorithms. SARIMA will handle daily and weekly cycles, while Prophet will model overarching trends and holiday effects with temperature as an exogenous variable [7]. To overcome traditional time series limitations in capturing non-linear weather interactions, tree-based models, specifically Random Forest, XGBoost, and LightGBM, will be applied using calendar features, seasonal terms and weather inputs [7,9-12,17]. Neural networks such as LSTM and support vector machines will remain secondary options due to excessive tuning overhead for structured tabular data [13-15]. Additionally, synthetic demand generation will be incorporated to model full distribution scenarios, directly addressing the reliance on single point forecasts identified in existing literature [7,8]. These modelling choices directly support the research question by enabling performance comparisons across seasonal variation, temperature changes and differing demand levels.

Models will be assessed using standard error metrics, including mean absolute error, root mean squared error and mean absolute percentage error, consistent with established practice in the forecasting literature [7,9–12]. Because the data are time-ordered, training will occur on earlier periods and testing on later periods, with a separate validation set used for tuning to avoid leakage. As we will be comparing multiple models, we will use consistent features, test and train sets, alongside standardised model evaluation (MSE, MAE and MAPE).

The quantitative workflow will rely primarily on Python libraries including pandas, NumPy, statsmodels, pmdarima, Prophet, scikit-learn, XGBoost, and LightGBM to maintain reproducible modelling pipelines [5,7]. TensorFlow and Keras will be reserved solely for potential LSTM implementations, while R tools such as ggplot2 and R Markdown will provide visualisation and reporting capabilities [5].

The project uses three NSW datasets from 2010 onwards, sourced from National Electricity Market and meteorological archives [1]. The primary target variable will be half-hourly electricity demand, paired with irregular temperature readings from the Bankstown weather station and historical operator forecasts. Data preparation will require standardising timestamps, resampling irregular temperature data to a matching half-hourly grid, imputing short gaps, and engineering calendar and lagged demand features.

This methodology assumes that seasonal patterns and temperature effects remain stable enough over the medium term to support model training, and that Bankstown temperature readings reasonably represent conditions across the NSW region. It also assumes that resampling irregular temperature data preserves the underlying weather signal. Finally, it assumes that operator forecasts included in the dataset reflect realistic baseline performance, allowing meaningful comparison between statistical and machine-learning approaches under varying seasonal and demand conditions.

4. Activities and Schedule

Group A meets three times weekly: a Thursday consultation with Wei Tian, a team stand-up immediately after, and a Sunday working session. Extra meetings occur before submissions. Swapnil Shinde, elected team leader in Week 0, chairs all sessions; agendas are posted to GitHub 24 hours before consultations and minutes within 24 hours after. The 33 activities sit on a Miro board maintained by Xavier Sun, each with a phase, owner, effort estimate, priority and assessment link. Code merges to the main branch only through reviewed pull requests, and an activity is complete once merged, rerun end-to-end from cleaned data, and written into the report.

The activities fall into six phases (Appendix B), each feeding the next. Work is front-loaded into data preparation because medium- to long-term forecasting depends more on design than fitting. Recent demand lags do not exist months ahead, so features are limited to calendar terms, harmonic seasonality, long-run trend and temperature. Temperature is also unknown at long horizons, so each model is fitted twice: once with observed temperature and once with analogue-year temperature, with the gap treated as a result. Modelling begins once the feature set and rolling-origin splits are finalised on 12 September. The six models run in parallel, meet a shared evaluation harness on 16 September, and freeze on 20 September ahead of the Week 4 checkpoint.

Appendix C table outlines each member's role. The structure is symmetric: every member owns one data or exploratory task, one forecasting model, and one report chapter. On 3 September the group agreed that each member should build and validate a model end-to-end so all can speak directly to the modelling process and avoid bottlenecks. Whoever drafted a section of this plan owns the corresponding report chapter, ensuring reading done now is not duplicated in Week 6. The board allocates 14–20 days of work per member across seven weeks, sized against availability recorded in Week 0.

5. References

- [1] Australian Energy Market Operator (AEMO). National Electricity Market (NEM). Available from: <https://www.aemo.com.au/energy-systems/electricity/national-electricity-market-nem>
- [2] Australian Energy Market Commission (AEMC). Congestion Management Review: Final Report, Appendix A – An introduction to congestion in the NEM. Sydney: AEMC; 2008.
- [3] Allen, C., Reid, M., Thwaites, J., Glover, R. & Kestin, T. (2020). Assessing national progress and priorities for the Sustainable Development Goals (SDGs): experience from Australia. *Sustainability Science*, 15(2), 521–538.
- [4] Hurn, S., Martin, V. & Tian, J. (2023). Modeling Multi-horizon Electricity Demand Forecasts in Australia: A Term Structure Approach. *The Energy Journal*, 44(3), 251–266. <https://doi.org/10.5547/01956574.44.2.shur>
- [5] Nunn, O. (2026). Project brief: Energy forecast [video lecture]. ZZSC9020 Data Science Project, UNSW Sydney. Available via UNSW Moodle.
- [6] Lafaye de Micheaux P, Drouilhet R, Liquet B. The r software: Fundamentals of programming and statistical analysis. Springer New York; 2013.
- [7] National Energy Retail Law (NSW) No 37a of 2012 (NSW), s 13, “National energy retail objective.” [Online]. Available: <https://legislation.nsw.gov.au/view/html/inforce/current/act-2012-37a#sec.13>. [Accessed: Sep. 3, 2026].
- [8] Husnoo, M. A., Anwar, A., Hosseinzadeh, N., Islam, S. N., Mahmood, A. N., & Doss, R. (2022). FedREP: Towards horizontal federated load forecasting for retail energy providers. *Proceedings of the 14th Asia-Pacific Power and Energy Engineering Conference (APPEEC)*, 1–6. IEEE. arXiv:2203.00219
- [9] Hussain, I., Ching, K. B., Uttraphan, C., Tay, K. G., & Noor, A. (2025). Evaluating machine learning algorithms for energy consumption prediction in electric vehicles: A comparative study. *Scientific reports*, 15(1), 16124.
- [10] Karunarathne, D. (2025). Machine Learning and Statistical Approaches for Forecasting Electricity Demand in Vaasa.
- [11] Hossain, S., Hasanuzzaman, M., Hossain, M., Amjad, M.H.H., Shovon, M.S.S., Hossain, M.S. and Rahman, M.K., 2025. Forecasting energy consumption trends with machine learning models for improved accuracy and resource management in the USA. *Journal of Business and Management Studies*, 7(1), pp.200-217.
- [12] Dudek, G. (2022). A comprehensive study of random forest for short-term load forecasting. *Energies*, 15(20), 7547.
- [13] Mazibuko, T., & Akindeji, K. (2025). Hybrid forecasting for energy consumption in South Africa: LSTM and XGBoost approach. *Energies*, 18(16), 4285.
- [14] Edoka, E. O., Abanihi, V. K., Amhenrior, H. E., Evbogbai, E. M. J., Bello, L. O., & Oisamoje, V. (2023). Time series forecasting of electrical energy consumption using deep learning algorithm. *Nigerian Journal of Technological Development*, 20(3).
- [15] Lian, H., Wei, H., Wang, X., Chen, F., Ji, Y., & Xie, J. (2025). Research on real-time energy consumption prediction method and characteristics of office buildings integrating occupancy and meteorological data. *Buildings*, 15(3), 404.
- [16] Fahim, K. E., De Silva, L. C., & Yassin, H. (2025, June). Time series forecasting of active power using ARIMA, SARIMA and hybrid models. In *Smart Sustainable Development Conference 2025 (SSD 2025)* (pp. 203-219). Atlantis

Press.

[17] Maleki, N., Lundström, O., Musaddiq, A., Jeansson, J., Olsson, T., & Ahlgren, F. (2024). Future energy insights: Time-series and deep learning models for city load forecasting. *Applied Energy*, 374, 124067.

[18] Jahangir, M.N. (2025). Machine Learning Techniques for Demand Forecasting in the Electricity Sector. In *Forecasting Methods for Renewable Power Generation* (eds J.G. Singh, R.K. Pachauri and S. Sreedharan).

<https://doi.org/10.1002/9781394249466.ch6>

[19] Kivimäki, M., Batty, G. D., Pentti, J., Suomi, J., Nyberg, S. T., Merikanto, J., ... & Vahtera, J. (2023). Climate change, summer temperature, and heat-related mortality in Finland: multicohort study with projections for a sustainable vs. fossil-fueled future to 2050. *Environmental health perspectives*, 131(12), 127020.

[20] Fitchett, J. M., Robinson, D., & Hoogendoorn, G. (2017). Climate suitability for tourism in South Africa. *Journal of Sustainable Tourism*, 25(6), 851-867.

[21] Liang, P., Yan, Z. W., & Li, Z. (2022). Climatic warming in Shanghai during 1873–2019 based on homogenised temperature records. *Advances in Climate Change Research*, 13(4), 496-506.

[22] Yusuf, N., Okoh, D., Musa, I., Adedola, S., & Said, R. (2017). A study of the surface air temperature variations in Nigeria. *The Open Atmospheric Science Journal*, 11(1), 54-70.

[23] Christy, J. R. (2026). Declines in hot and cold daily temperature extremes in the conterminous US, 1899–2025. *Theoretical and Applied Climatology*, 157(5), 309.

[24] Hasan, D. S. N. A. B. P. A., Ratnayake, U., & Shams, S. (2016, February). Evaluation of rainfall and temperature trends in Brunei Darussalam. In *AIP Conference Proceedings* (Vol. 1705, No. 1, p. 020034). AIP Publishing LLC.

[25] Limjirakan, S., & Limsakul, A. (2012). Observed trends in surface air temperatures and their extremes in Thailand from 1970 to 2009. *Journal of the Meteorological Society of Japan. Ser. II*, 90(5), 647-662.

6. Appendix

Appendix A — Reported Model Performance Metrics

Table A1. Electricity demand, Vaasa — short-term hourly [10]

Model	MSE	RMSE	MAE	MAPE	R ²
XGBoost	21.22	4.61	3.61	0.063	0.902
Multiple Linear Regression	21.40	4.626	3.66	0.066	0.901
Random Forest	23.19	4.82	3.67	0.061	0.892

Table A2. U.S. energy consumption trends — classification metrics [11]

Model	Accuracy	Precision	Recall	F1
Logistic Regression	0.701	0.70	0.70	0.70
Random Forest	0.696	0.70	0.70	0.70
XGBoost	0.67	0.67	0.67	0.67

Table A3. South African energy consumption [13]

Model	MAE
LSTM	0.2
XGBoost	—
ARIMA	—

Table A4. Transmission Company of Nigeria, Aug–Dec 2021 — half-hourly load [14]

Model	MAPE	RMSE
LSTM	0.01	19.759

Table A5. Shanghai office building, 24 Jun – 15 Sep 2013 [15]

Model	MAPE
LSTM	9%
Backpropagation Network (BP)	12%
Artificial Neural Network (ANN)	16%
Support Vector Machine (SVM)	21%

Table A6. LSTM timestep optimisation, Shanghai office building [15]

Timestep	MAPE
----------	------

6 h	3%
18 h	5%
24 h	6%

Table A7. Active power usage, Brunei — short-term [16]

Model	RMSE
ARIMA	0.048
SARIMA	0.6201
Holt-Winters + SARIMA (hybrid)	1.15

Table A8. City energy consumption, Kalmar, Sweden — training [17]

Model	MSE	RMSE	MAE	MAPE	R ²
Random Forest	64.61	8.03	6.19	4.67%	66%
XGBoost	119.82	10.94	8.63	6.51%	38%
SARIMAX	91.29	9.55	3.71	2.75%	52%
FB Prophet	—	—	—	—	—
CNN	13.36	8.21	2.75	0.02%	0.93
ARIMAX	—	—	—	—	—

Table A9. City energy consumption, Kalmar, Sweden — testing [17]

Model	MSE	RMSE	MAE	MAPE	R ²
Random Forest	278.19	16.67	14.43	9.35%	64%
XGBoost	—	—	—	—	—
SARIMAX	133.72	11.56	9.87	6.28%	20%
FB Prophet	97.43	9.87	7.95	5.22%	0.42
CNN	67.60	8.22	5.74	4.35%	0.60
ARIMAX	—	—	—	—	—

Dataset: hourly datapoints 2021–2023, features include temperature and price; FB Prophet additionally models holidays.

Table A10. Temperature Ranges (Sweden; USA; South Africa; Nigeria; Brunei; Shanghai) [19-25]

Location	Temperature Range
Sweden	~29°C
USA	~22°C

South Africa	~32°C
Nigeria	~32°C
Brunei	~10°C
Shanghai	~33°C
Thailand	~22°C

Appendix B — Phases, activities and outputs

Phase	Activities (Task ID)	Est. Date Range	Planned Output
Project Planning	Task 1 to 8; Task 33	31/08/2026 ~ 08/09/2026	Working contract, agreed research question, the project plan and each member's checklist
Data Engineering & Preparation	Task 10 to 11	07/09/2026 ~ 12/09/2026	Cleaned half-hourly dataset, horizon-safe feature set, rolling - origin splits
Exploratory Data Analysis	Task 12 to 13	09/09/2026 ~ 12/09/2026	Load profiles, the temperature-demand response curve, trend and structural breaks
Modelling	Task 14 to 19	14/09/2026 ~ 15/09/2026	Six forecasters built in parallel, one per member, on a common feature set
Evaluation & Validation	Task 20 to 20	16/09/2026 ~ 23/09/2026	Stratified error tables, a frozen model set, one improvement round on the winner
Reporting & Communication	Task 9; Task 23 to 32	21/09/2026 ~ 13/10/2026	Weekly reflections; the final 35-page group report and 15-minutes video.

Appendix C — Roles and justification

Member	Role	Assignment Basis
Swapnil Shinde	Team Lead	Elected in Week 0, and holds the broadest skill profile in the group: Python, R, SQL, statistics, visualisation and reporting. Chairs the stand-up and the lecturer consultation, arbitrates where the group is split, and holds the statistical stream, the temperature–demand response (Task 13) followed directly by SARIMAX (Task 15). Writes up software, data and pre-processing (Task 26).
Daniel Cao	Machine Learning Engineer	ML engineering and data science experience, and custodian of the group repository. Owns the ingest-and-clean pipeline (Task 10) and the Random Forest (Task 16), then writes the data cleaning, assumptions and modelling methods (Task 27) documenting work done firsthand rather than inherited.
Kelly Hamilton	Data Science Research	A psychology background trained in research design and literature synthesis, with statistics and machine learning on top. Owns the literature review and the discussion (Task 25), the two chapters that frame the project and then interpret its results, and builds the LSTM (Task 19) so that the deep learning comparison is argued by the member writing the discussion.
Faiza Mumtaz	Data Scientist	ML engineering combined with data analysis, Tableau and reporting. Produces the load-profile and seasonality analysis (Task 12) and the XGBoost model (Task 17), then writes the Exploratory Data Analysis chapter (Task 28) built from that same analysis.
Arman Hajisafi	Data Engineer	PySpark and data engineering experience, the direct match for the horizon-safe feature set and the rolling-origin split design (Task 11) over eleven years of half-hourly data, the single activity every later result depends on. Also builds LightGBM (Task 18), pairing with XGBoost for a like-for-like comparison, and writes the introduction, conclusion and reference organisation (Task 24).
Xavier Sun	Product	Works as a product manager for data science and ML engineering teams, so owns the board, schedule and risk log (Task 33), and keeps the plan current as scope moves. A commerce and accounting background anchors the translation of model output into the business framing the non-technical needs. Builds the baseline models (Task 14) that every other result is measured against, writes the analysis and results chapter (Task 29), and produces the video (Task 32).

Appendix D - DS Project Task Planner (Group A)

Task Board												
	Tasks	Phase	Assignee	Status	#	LOE (D)	Priority	Description	Start Date	End Date	Assessment	
1	Team Contract, roles, tooling setup and finalise Research Topic	Project Planning	All Members	Complete	1		High	Agree the working contract and decision rules, confirm roles, and stand up Teams, Planner, OneDrive and the GitHub repository.	Sep 1, 2026	Sep 3, 2026	A1A - Plan	
2	Weekly Lecturer Consultation	Project Planning	All Members	In Progress	7		Medium	Book the slot, circulate an agenda to the repository 24 hours before, publish minutes within 24 hours after.	Sep 3, 2026	Oct 16, 2026	All	
3	Project Plan S1: Introduction and Motivation	Project Planning	Daniel Cao	Complete	2		Medium	Write section 1: the business problem, why medium-term demand forecasting matters, who the client is, and what the project sets out to achieve.	Sep 1, 2026	Sep 6, 2026	A1A - Plan	
4	Project Plan S2: Brief Literature Review	Project Planning	Arman Hajisafi Kelly Hamilton	Complete	4		High	Write section 2: survey existing work on medium- and long-term load forecasting, identify the methodological gap, and state how this project addresses it.	Sep 1, 2026	Sep 6, 2026	A1A - Plan	
5	Project Plan S3: Methods, Software and Data Description	Project Planning	Faiza Mumtaz Swapnil Shinde	Complete	4		High	Write section 3: the intended models & why they suit the question, software stack, and a description of the three datasets including format, size and known quality issues.	Sep 1, 2026	Sep 6, 2026	A1A - Plan	
6	Project Plan S4: Activities and Schedule	Project Planning	Xavier Sun	Complete	2		Medium	Write section 4: the activity list, the phase structure, role justification, the timeline and the Gantt chart.	Sep 1, 2026	Sep 6, 2026	A1A - Plan	
7	Group Review Draft - Project Plan	Project Planning	All Members	Complete	1		Medium	Read the four sections together, resolve overlaps and contradictions, check the three-page limit, and submit.	Sep 6, 2026	Sep 6, 2026	A1A - Plan	
8	Complete Individual Implementation Checklist	Project Planning	All Members	Complete	1		High	Assessment 1B: Each member's checklist by 5 pm Tuesday of Week 2.	Sep 5, 2026	Sep 8, 2026	A1B - Checklist	
9	Complete Individual Weekly Self-Reflection	Reporting & Communication	All Members	In Progress	4		Medium	An individual task: write a 500-word summary reflective portfolio, based on your weekly teamwork reflections for Weeks 1-4.	Aug 31, 2026	Sep 29, 2026	A2 - Reflection	
10	Ingest validate and clean datasets	Data Engineering & Preparation	Daniel Cao Kelly Hamilton	Not Started	2		Medium	Load the three CSVs and confirm columns, types, units, date coverage and row counts against the course brief, check duplicates, and resample temperature onto the 30-minute demand grid. External Data: @Kelly to source Solar-Radiation dataset and dataset clean-up.	Sep 7, 2026	Sep 9, 2026	A3 - Report	
11	Build feature set and design rolling-origin splits	Data Engineering & Preparation	Arman Hajisafi	Not Started	3		Medium	Calendar and public + school-holiday flags, Fourier seasonality terms, a long-run trend and cooling/heating degree days on a fixed 18 C base, which activity 13 then validates rather than supplies; set the train and test origins for each horizon, then review every feature for information that would not exist at forecast time.	Sep 9, 2026	Sep 12, 2026	A3 - Report	
12	Load profiles and seasonality	Exploratory Data Analysis	Faiza Mumtaz	Not Started	2		Medium	Demand by hour, weekday, month and year; STL decomposition and autocorrelation to size the seasonal structure the models have to capture.	Sep 9, 2026	Sep 11, 2026	A3 - Report	
13	Temperature-demand response and trend and structural breaks	Exploratory Data Analysis	Swapnil Shinde	Not Started	3		Medium	Fit the U-shaped response curve, estimate the heating and cooling thresholds, and quantify how many megawatts each additional degree costs. Quantify the daytime decline driven by rooftop solar across 2010-2021 (isolate the 2020 COVID-19 period).	Sep 9, 2026	Sep 12, 2026	A3 - Report	
14	Baseline Modelling	Modelling	Xavier Sun	Not Started	3		High	Seasonal naive (same period last year) and a climatology profile averaged over prior years. These set the bar every other model has to clear, and every later result is reported relative to them.	Sep 12, 2026	Sep 15, 2026	A3 - Report	
15	Statistical models - SARIMAX	Modelling	Swapnil Shinde	Not Started	3		High	Fit SARIMAX with temperature as an exogenous input, with ETS and Prophet as comparators; check residual diagnostics and report the fitted seasonal structure.	Sep 12, 2026	Sep 15, 2026	A3 - Report	
16	ML Modelling - Random Forest	Modelling	Daniel Cao	Not Started	3		High	Train a random forest on the horizon-safe feature set, tuned only on validation origins so the test origins stay untouched; record feature importances for the discussion.	Sep 12, 2026	Sep 15, 2026	A3 - Report	
17	ML Modelling - XGBoost	Modelling	Faiza Mumtaz	Not Started	3		High	Train a gradient-boosted tree ensemble with XGBoost on the same features and splits; tune depth, learning rate and regularisation on validation origins only.	Sep 12, 2026	Sep 15, 2026	A3 - Report	
18	ML Modelling - LightGBM	Modelling	Arman Hajisafi	Not Started	3		High	Train a LightGBM ensemble on the same features and splits, and compare its training cost and accuracy against XGBoost, which uses the same algorithm family with a different split-finding strategy.	Sep 12, 2026	Sep 15, 2026	A3 - Report	
19	ML Modelling - LSTM	Modelling	Kelly Hamilton	Not Started	3		High	Train an LSTM sequence model on the same features and splits, and report whether the extra capacity earns its cost at these horizons.	Sep 12, 2026	Sep 15, 2026	A3 - Report	
20	Model evaluation & validation	Evaluation & Validation	All Members	Not Started	2		Medium	Each member puts their own model through the shared harness: fit at every rolling origin, predict the horizon, collect MAE, RMSE and MAPE. Then break the error down by season, temperature band and demand level, test summer peaks and extreme-heat days.	Sep 16, 2026	Sep 18, 2026	A3 - Report	
21	Model selection and freeze	Evaluation & Validation	All Members	Not Started	1		Medium	Compare the six models on the common metrics, choose the best per horizon, record the decision and the evidence behind it, and stop experimenting.	Sep 19, 2026	Sep 20, 2026	A3 - Report	
22	Selected model improvement	Evaluation & Validation	Model Winner	Not Started	2		Medium	The member whose model wins takes one focused round of tuning or feature refinement, with the rest of the team reviewing. The owner is decided at Task 21 (Model Selection).	Sep 21, 2026	Sep 23, 2026	A3 - Report	
23	Report skeleton and section owners, and draft the report chapters	Reporting & Communication	All Members	Not Started	1		Medium	Agree headings, page budget and a named owner for each of the seven report chapters, then start drafting against the frozen results.	Sep 21, 2026	Sep 22, 2026	A3 - Report	
24	Group Report: Introduction + Conclusion and Further Issues + Organising Reference & Appendix	Reporting & Communication	Arman Hajisafi	Not Started	5		Medium	- Chapter 1: Introduction - Chapter 7: Conclusion and Further Issues - Reference + Appendix: Sorting & Organising	Sep 23, 2026	Oct 7, 2026	A3 - Report	
25	Group Report: Literature Review + Discussion	Reporting & Communication	Kelly Hamilton	Not Started	5		High	- Chapter 2: Literature Review - Chapter 6: Discussion	Sep 23, 2026	Sep 29, 2026	A3 - Report	
26	Group Report: Material and Methods: 3.1, 3.2 and 3.3	Reporting & Communication	Swapnil Shinde	Not Started	5		Medium	- Chapter 3: 3.1 Software 3.2 Description of the Data 3.3 Pre-processing Steps	Sep 23, 2026	Sep 29, 2026	A3 - Report	
27	Group Report: Material and Methods: 3.4, 3.5 and 3.6	Reporting & Communication	Daniel Cao	Not Started	5		High	- Chapter 3: 3.4 Data Cleaning 3.5 Assumptions 3.6 Modelling Methods	Sep 23, 2026	Sep 29, 2026	A3 - Report	
28	Group Report: Exploratory Data Analysis	Reporting & Communication	Faiza Mumtaz	Not Started	5		Medium	Chapter 4, written from the EDA notebooks, with every figure regenerated from the cleaned data rather than pasted.	Sep 23, 2026	Sep 29, 2026	A3 - Report	
29	Group Report: Analysis and Results	Reporting & Communication	Xavier Sun	Not Started	5		High	Chapter 5: the model comparison, the stratified error tables, and what they say about which model wins under which conditions.	Sep 26, 2026	Oct 2, 2026	A3 - Report	
30	Internal review and reproducibility check, and final submission	Reporting & Communication	All Members	Not Started	4		Medium	Clone the repository fresh, re-run the pipeline end to end, confirm every figure and number in the report regenerates, then submit.	Sep 29, 2026	Oct 7, 2026	A3 - Report	
31	Video storyboard and slides	Reporting & Communication	All Members	Not Started	2		Low	Storyboard fifteen minutes for a non-technical audience and build one shared deck rather than six.	Oct 7, 2026	Oct 10, 2026	A4 - Presentation	
32	Record, edit and submit the video	Reporting & Communication	Xavier Sun	Not Started	2		Medium	Each member records their own section; edit the pieces together and submit by 10 pm Tuesday of Week 7.	Oct 10, 2026	Oct 13, 2026	A4 - Presentation	
33	Project coordination	Project Planning	Xavier Sun	In Progress	4		Medium	Maintain the task board and Gantt, keep the risk log current, write the weekly agendas and minutes, and enforce the branch-and-review convention on the repository.	Aug 31, 2026	Oct 13, 2026	All	

+

Appendix E - DS Project Task Planner Timeline - Gantt (Group A)

